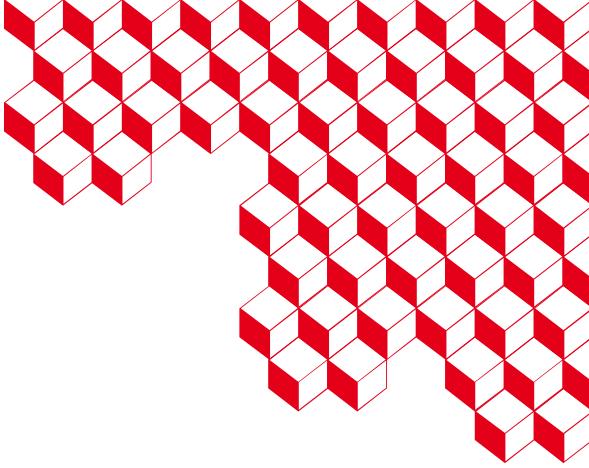




list



Symbolic Knowledge Injection in LLMs for Zero-Shot and Few-Shot Scenarios

Intern: Valentin Six

Advisors: Gaël de Chalendar, Evan Dufraisse

Table of content

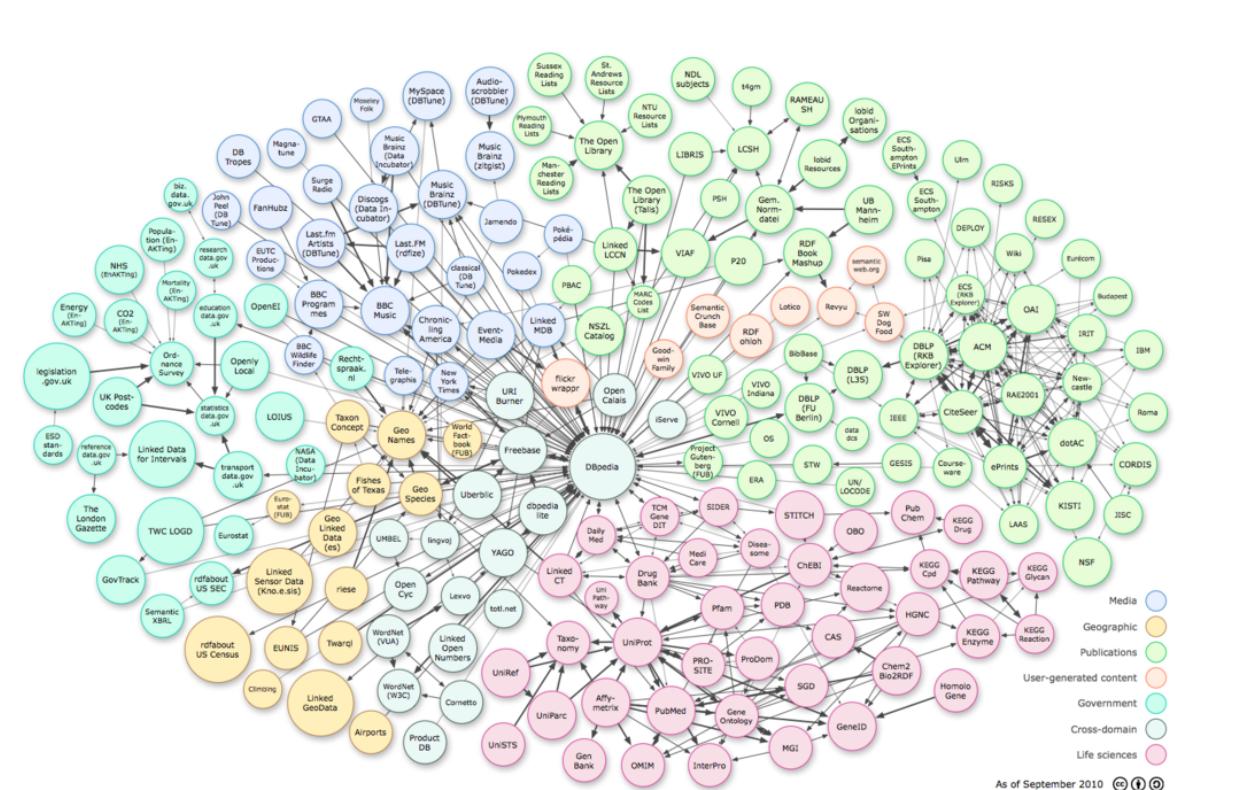
1. Research subject
2. Knowledge Injection
3. Reasoning Mechanisms
4. Contributions
5. Detailed Method
6. References



Research Subject

- LLMs can lack explicit understanding and reasoning capabilities
- For certain tasks, data can be rare (or absent): training becomes difficult or expensive...

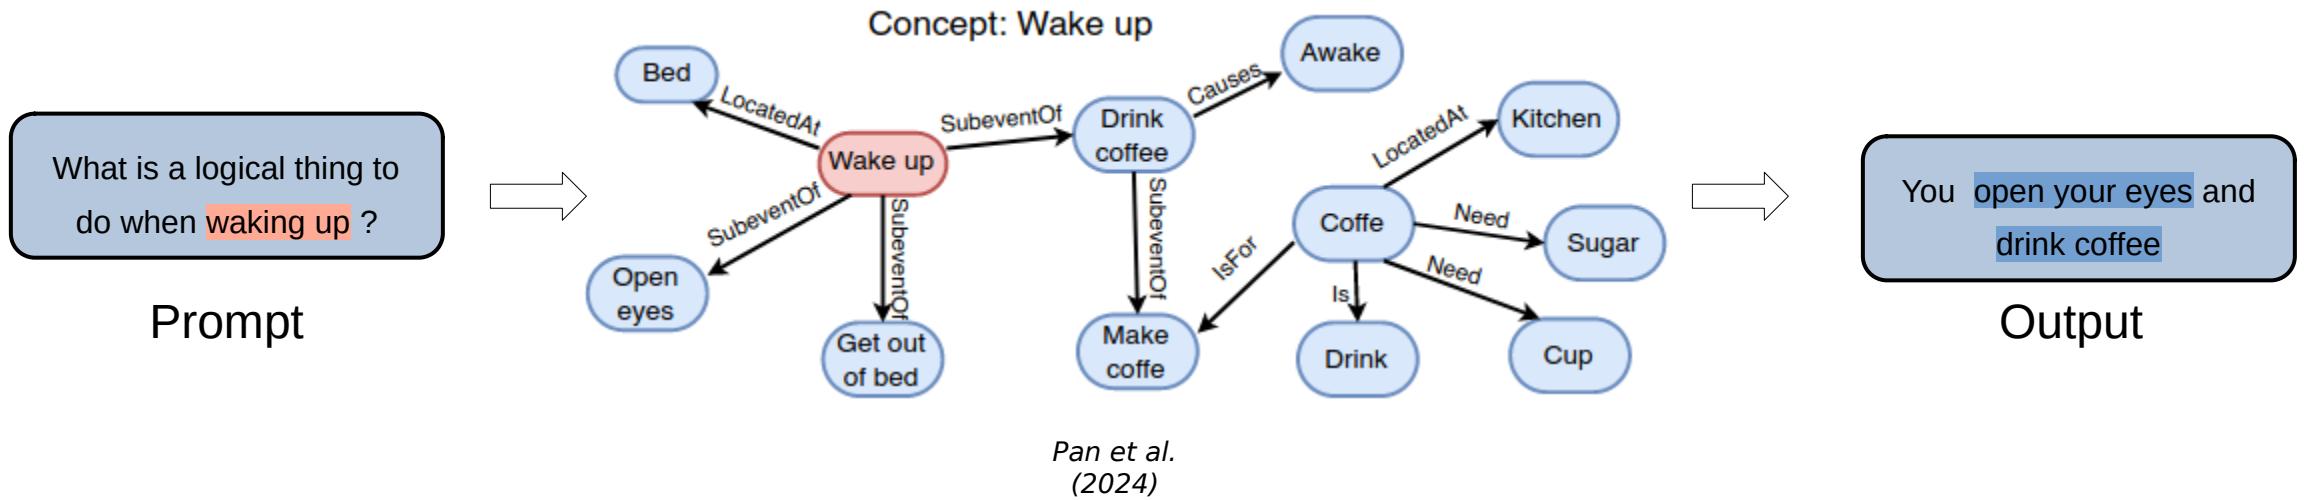
→ Objective: use **knowledge graphs** and exploit richness of **structure** to improve reasoning capacity of LLMs + ground the model with facts



<https://accidental-taxonomist.blogspot.com/2019/05/knowledge-graphs-and-ontologies.html>



Knowledge Injection



- Enable better **question understanding + reduce hallucinations**
- Inject facts / implicit relationships (not mentioned in the prompt)



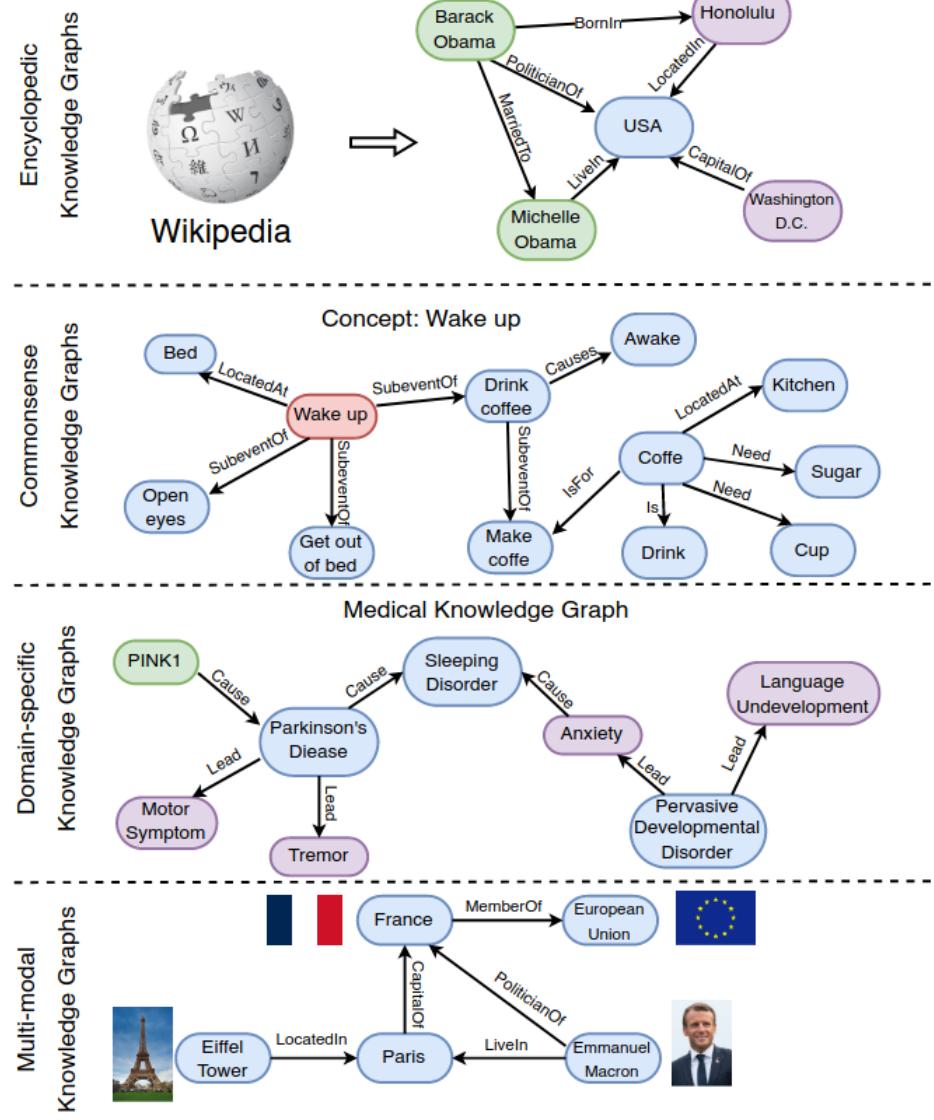
Knowledge Injection

Why using **knowledge graphs** ?

- Density of information + rich structure
- Traceability + results interpretation
- Different graph types for different use cases

How to use them ?

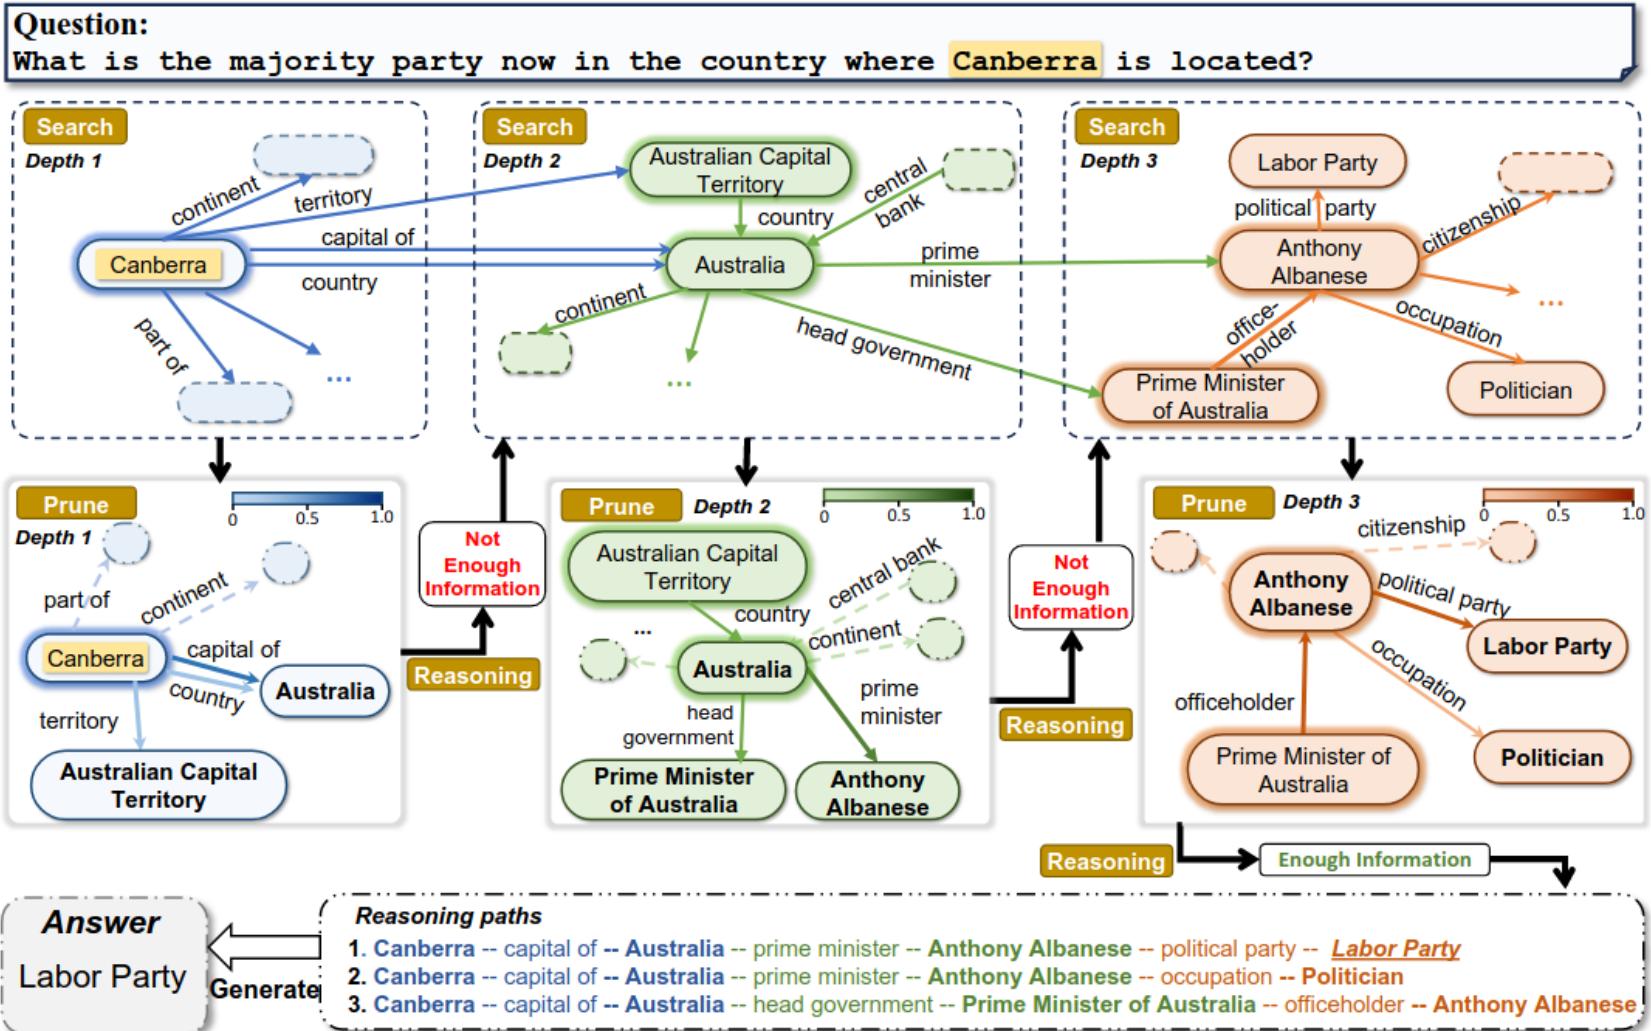
- Triplets, reasoning paths, sub-graphs
- Methods : GraphRAG, fine-tuning, pre-training



Pan et al.
(2024)



How about reasoning ?





How to improve reasoning ?

Graph retrieval is a difficult task + expensive task for complex queries

Intuition :

- **Decomposing** complex questions in simple sub-questions and perform multiple retrievals
- Fully capture **all dimensions** of the complex question
- Use **reasoning** used within decomposition to improve graph retrieval

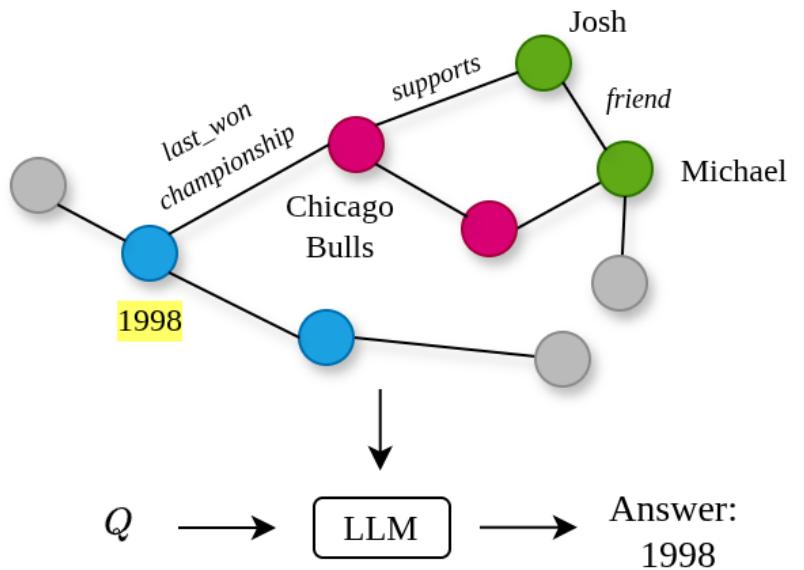
Using “G-Retriever” approach to handle textual graphs.

Q : When did the team that Michael's best friend support last win the Championship ?

q_1 : Who is Michael's best friend ?

q_2 : What team does he support ?

q_3 : When did that team last win the Championship ?





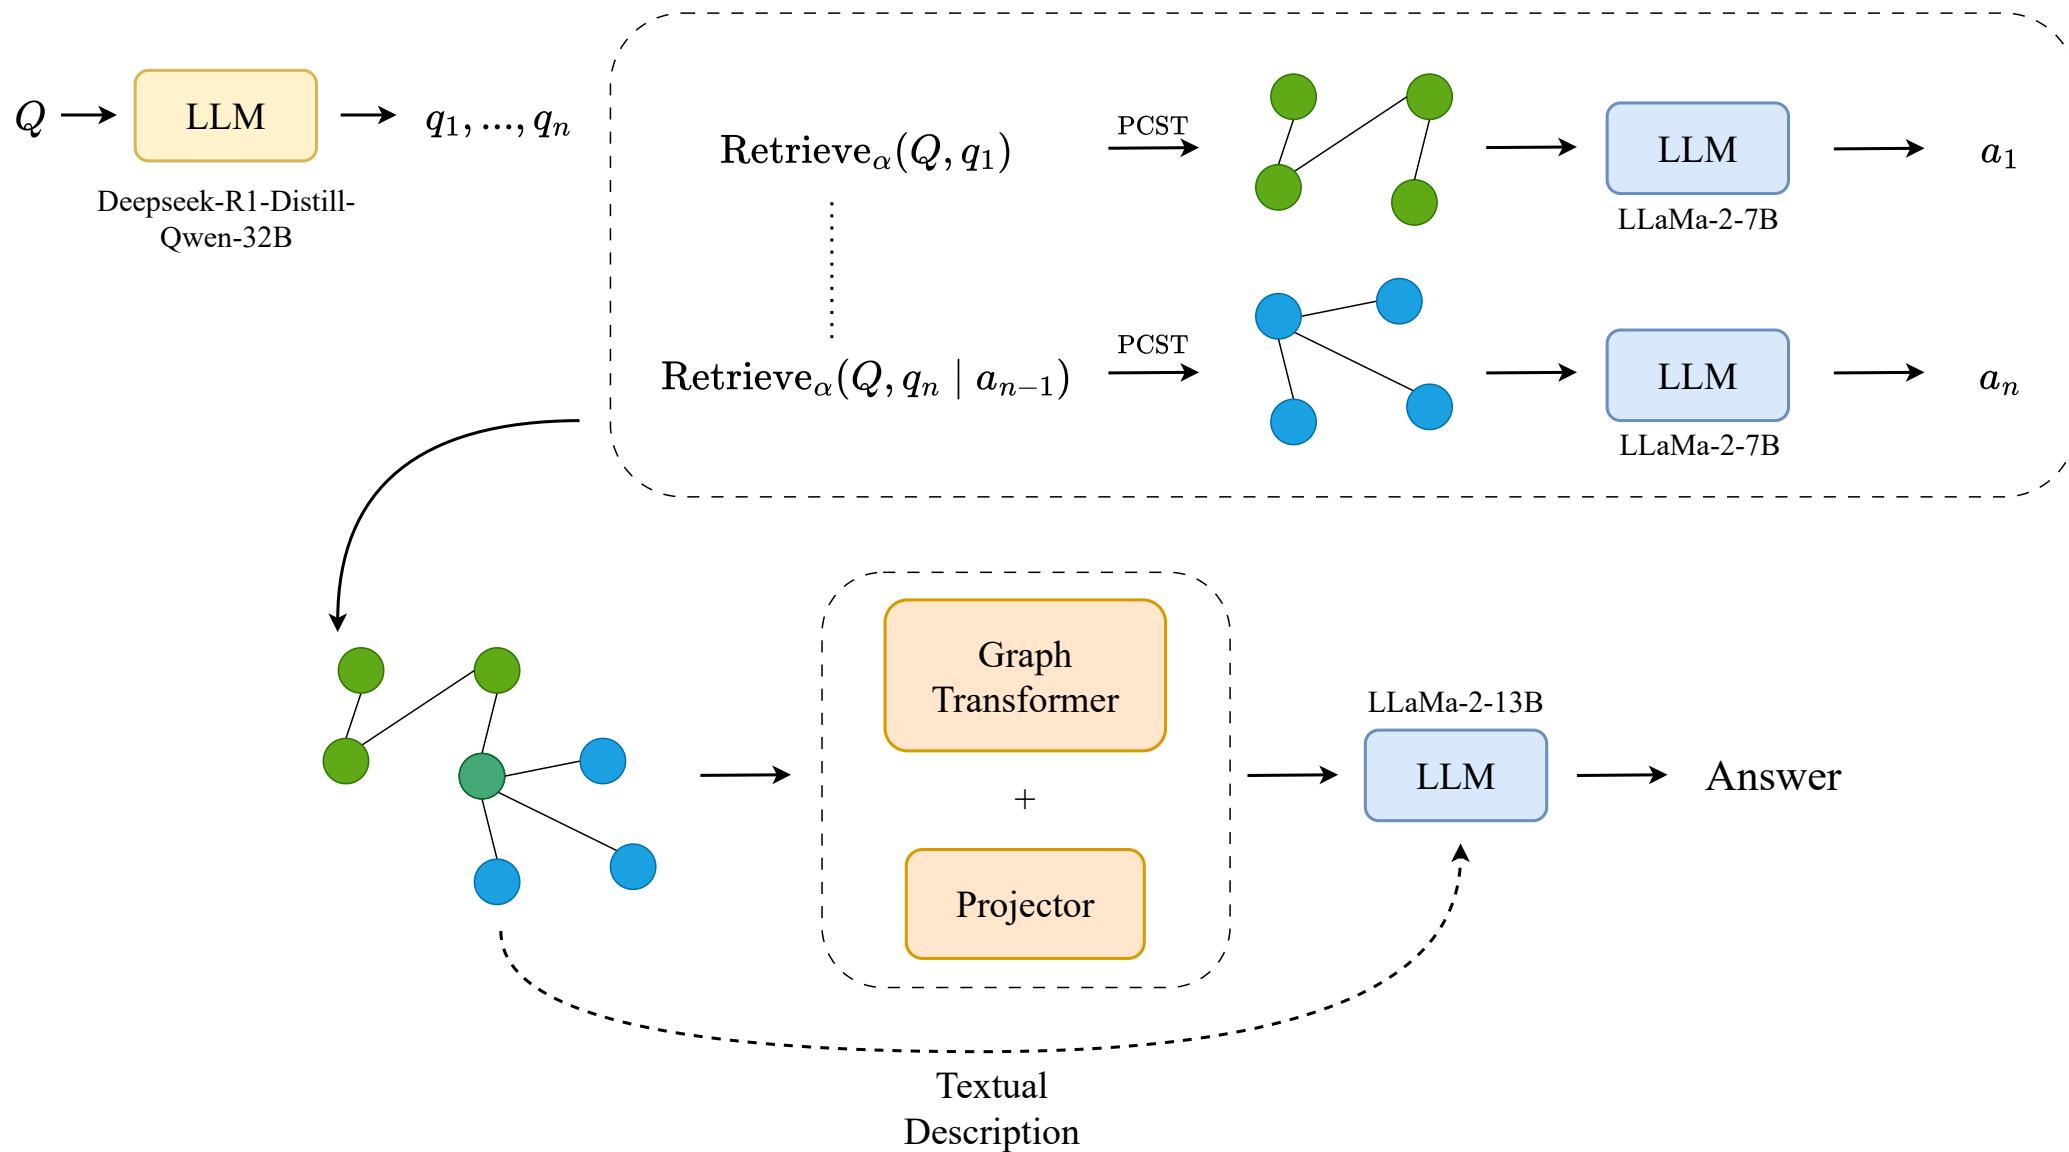
WebQSP vs. CWQ Benchmarks

- **WebQSP**: QA dataset, based on Freebase ; simple questions (2-hop max)
→ *used by G-Retriever*
- **CWQ** : also based on Freebase ; more complex questions (multi-hop)
→ *ignored by G-Retriever ; would probably struggle*

Objective: answer more complex questions that require multiple steps + reasoning



Overall Generation Pipeline





Retrieval and Graph Construction

1) Knowledge retrieval for each sub-question

$$\left. \begin{array}{l} V_k^i = \underset{n \in V}{\operatorname{argtopk}} [\alpha \cos(z_{q_i}, z_n) + (1-\alpha) \cos(z_q, z_n)] \\ E_k^i = \underset{e \in E}{\operatorname{argtopk}} [\alpha \cos(z_{q_i}, z_e) + (1-\alpha) \cos(z_q, z_e)] \end{array} \right\} \text{Construct sub-graph } G_i = (V_k^i, E_k^i)$$

Problem: sub-questions lack self-awareness ; not capturing the “global objective”

- Retrieve the k most similar nodes + edges for each sub-question
- Using the initial question Q and current sub-question q_i

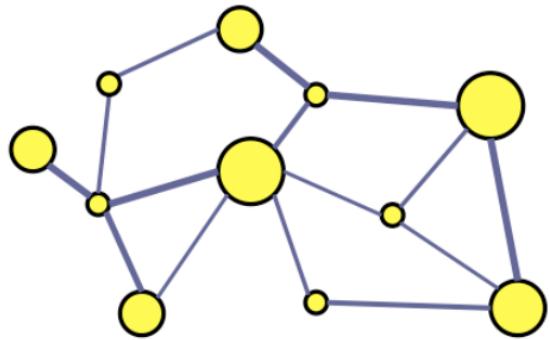
Parameter α to control grounding ; “double cosine similarity”



Retrieval and Graph Construction

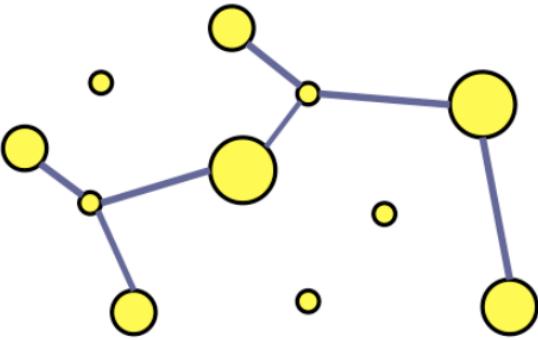
2) Graph transformation for each sub-question

→ Prize-Collecting Steiner Tree Algorithm (*TL;DR : most useful graph with minimal size*)



$$G_i = (V_k^i, E_k^i)$$

Retrieved Graph



$$G_i^* = PCST(V_k^i, E_k^i)$$

*Filtered + Connected
Graph*

Akhmedov et al.
(2018)



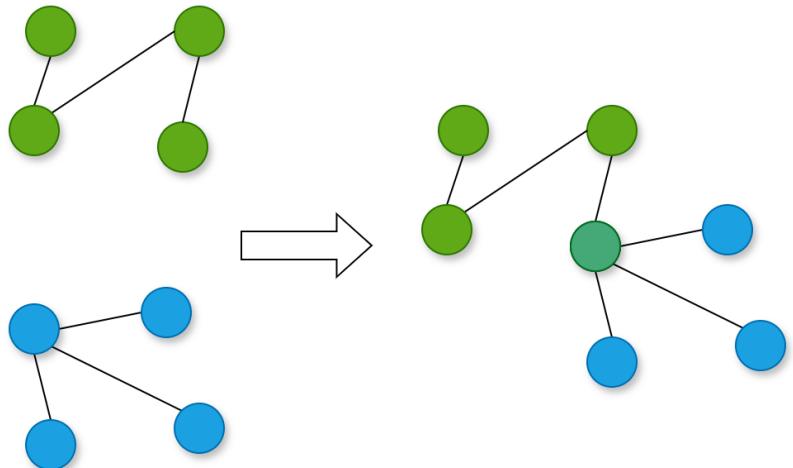
Retrieval and Graph Construction

3) Graphs merging

$$G_i^* = PCST(V_k^i, E_k^i) \longrightarrow G^* = \coprod_i^n (G_1^*, \dots, G_n^*)$$

Obtaining connected sub-graphs

Merging

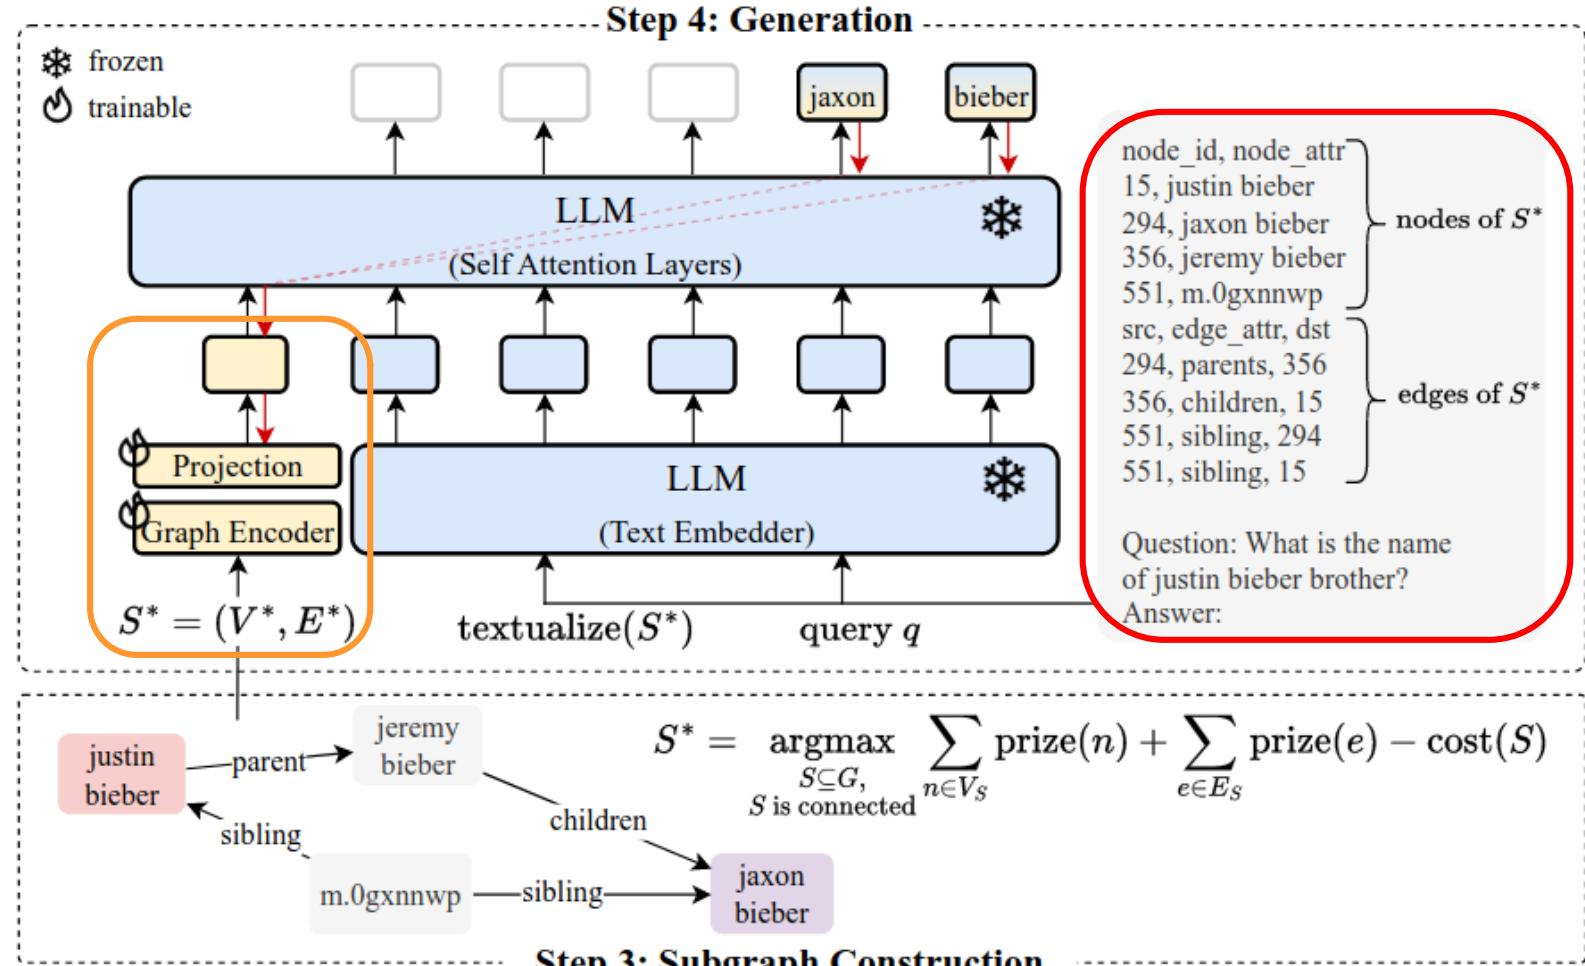


- For each question, we **merge** the sub-graphs corresponding to the sub-questions
- Detailed merging process:
 - Extract sets of unique nodes + edges
 - Create merged graph

Note: there is no guaranty that the final graph is connected, but most likely it will be



Answer Generation



Generation:

Hard prompt + Soft prompt



Reconstruction:

Connected sub-graph from retrieved elements

He et al. (2024)



Results

Method	CWQ		WebQSP	
	Hit@1	F1	Hit@1	F1
IO prompt (ChatGPT)	37.6	-	63.3	-
CoT (ChatGPT)	38.8	-	62.2	-
StructGPT (ChatGPT)	54.3	-	72.6	-
ToG (LLaMa-2-70B)	53.6	-	63.7	-
ToG (ChatGPT)	57.1	-	76.2	-
RoG (LLaMa-2-7B + FT)	<u>62.6</u>	56.2	85.7	70.8
PoG (GPT-3.5)	63.2	-	<u>82</u>	-
G-R (LLaMa-2-7B)	52.1	44.8	70.5	51.7
Ours (LLaMa-2-7B)	54.9	46	71.9	52.4
G-R (LLaMa-2-13B)	54.6	46.9	76.5	<u>57.2</u>
Ours (Hybrid 7B/13B)	<u>57.9</u>	<u>50.3</u>	77.9	58.2
Ours (LLaMa-2-13B)	58.1	50.8	<u>77.4</u>	56.4

Accuracy results compared to SOTA

Smaller models

Larger models \neq better performance

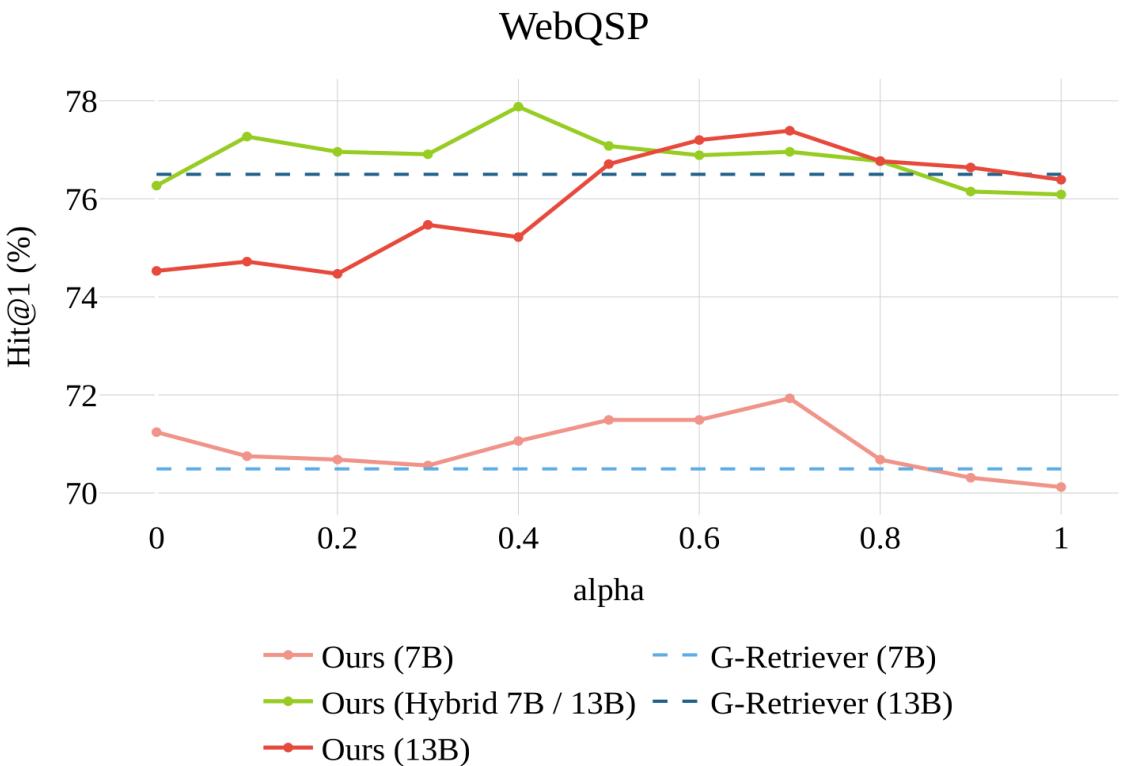
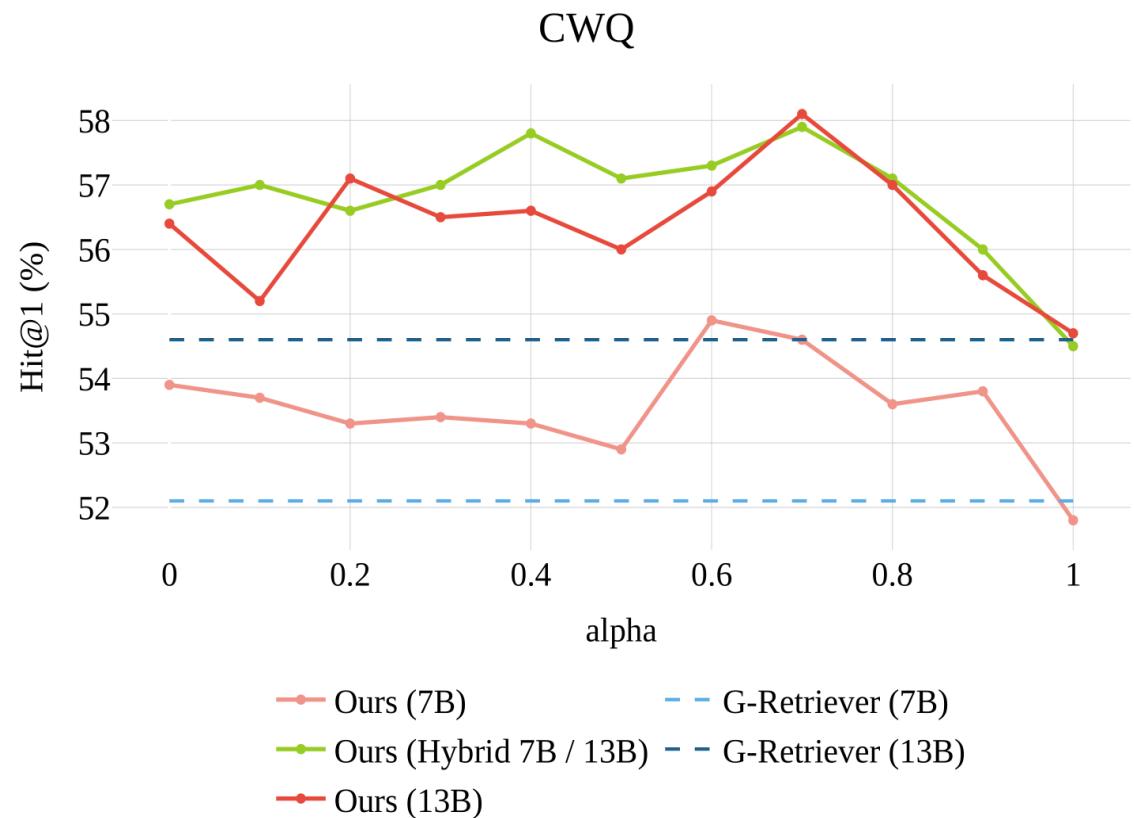
Fewer LLM calls

Method	CWQ	WebQSP
ToG	22.6	15.9
PoG	13.3	9.0
Ours	4.8	4.3

Number of LLM calls



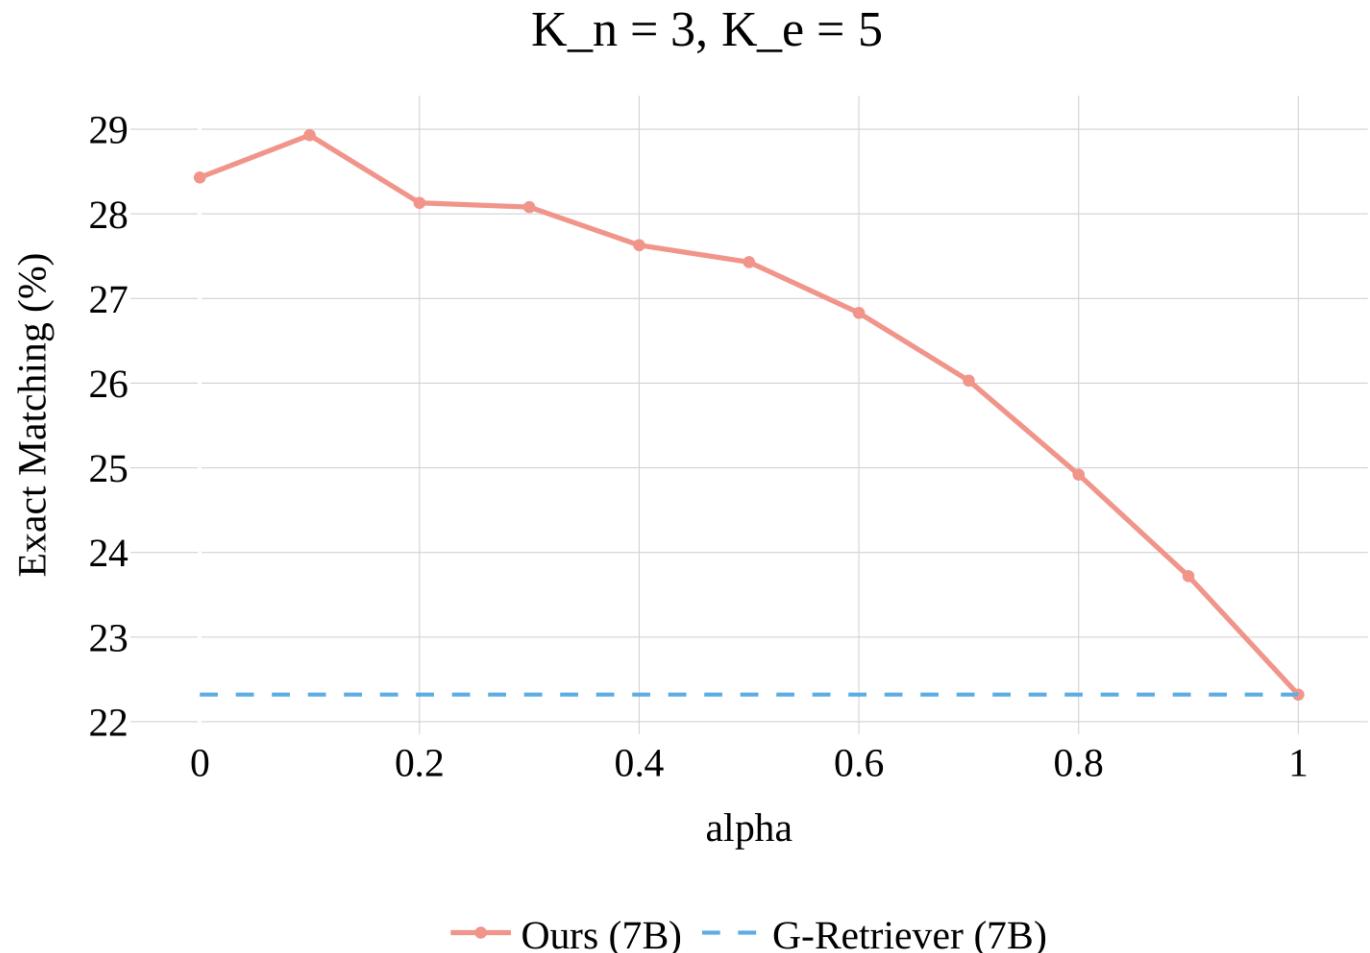
Results



Our retrieval methods improves the performance of G-Retriever for different model sizes. We introduce an hybrid approach where we alternate between 7B and 13B models; this approach achieves similar performances to the 13B model for the entire pipeline.



Results



In this context, Exact Matching corresponds to the percentage of retrieved graphs that contain the ground truth answer.

Focusing on the sub-questions gives better EM, at the cost of having disconnected and less dense graphs...



References

- « Unifying Large Language Models and Knowledge Graphs: A Roadmap » (Pan et. al)
- « Think-on-Graph: Deep and Responsible Reasoning of Large Language Model on Knowledge Graph » (Sun et al.)
- « G-Retriever: Retrieval-Augmented Generation for Textual Graph Understanding and Question Answering » (He et al.)